

Predicting Query Execution Time

Name

Mid-term ME Project Report

Abstract

The ability to estimate the query execution time is central for a number of tasks in database system such as query scheduling, progress monitoring and costing during query optimization. Recent work has explored the use of statistical techniques in place of the manually constructed cost models used in query optimization. Such techniques, which require as training data along with the actual execution time, promises superior accuracies for they being able to account the for factors such as hardware characteristics and bias in cardinality estimates. However, such techniques fail to generalize i.e., produce poor estimates for queries that are not seen during the training.

In this work, we propose and evaluate predictive modeling techniques that learn query execution behavior at a fine grained operator level. For each operator, we consider different sets of features and build different models for them. Since there are only finitely many operators in database, this approach is practical and will be able to estimate any query as its a composition of many operators. We evaluate our approaches using TPC-H and TPC-DS workloads on PostgreSQL.

1 INTRODUCTION

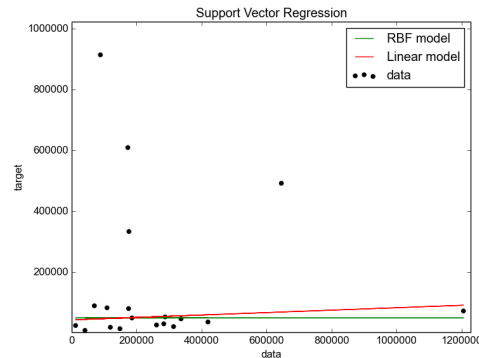
Database systems can greatly benefit from accurate execution time predictions including:

- Admission control: Resource managers can use this metric to perform workload allocations such that the specific QoS are met [?].
- Query Scheduling: Knowing the execution time is crucial in deadline and latency aware scheduling .
- Progress monitoring:

Currently, resource estimation is based on manually constructed models, which are part of the query optimizer and typically use combinations of weighted estimates of the number of tuples flowing through operators, column widths, etc. Unfortunately, such models often fail to capture several factors that affect the actual resource consumption. For example, they may not include detailed modeling of all of the various improvements made to database query processing such as nested loop optimizations [13, 11] which localize references in the inner subtree, and introduce partially blocking batch sorts on the outer side, thereby increasing the memory requirements and CPU time

and reducing I/O compared to the traditional iterator model. Similarly, they may not accurately reflect specific hardware characteristics of the current production system or the impact of cardinality estimation errors

If possible plot a non-linear SVR for the data. will the optimizer's cost itself is enough 1D regression using linear, polynomial and RBF kernels



2 Overview

Elaborate what is the plan level approach. give a intro model based techniques. i.e., OFF LINE TRAINING,

emphasize one time. 1. How to obtain training data (explain analyze) 2. getting execution times at plan nodes 3. Individual model training 4. Explain while testing and predict using Model more in next section

Mention assumptions, Pure in-memory.. lets solve this first. For example, a sort merge

join may require sorting its inputs. The sort operators are then responsible for handing over their result to the merge join via main memory. This means that the merge join may not require any I/O if the merge can be done purely in main memory

Analytical approaches problems:

counting the numbers of pages read is not sufficient as the discrepancy between random and sequential I/O is tremendous.

3 Model selection and Training

What model is considered good ? choice of model and kernel If possible prove linear systems are not sufficient. Training data ,specification. Does more training data helps? training time etc.. Essentially all operators need to be covered. generation QGEN tool.

4 Preliminary Experiments

Give a background about Static and dynamic workload. do experiments at 2 stages 1. Act vs Act 2. Est vs Est(Bias towards estimated cardinalities) For each type of above scenario, do testing for tpc-h queries (tpc-ds in future) At different scales to see the model's ability to predict. For midterm you should be having result for 1 as well as 10 GB. For each above dataset, Results should convey L1 (MRE) as well as queries that are within [0-0.5] [0.5-1] [1-1.5]

5 Conclusions and Future Work

Talk about generalization. operator specific features capturing query interaction i.e., pipeline source of over-estimation Optimizers provides sufficient information to detect a pipeline, can we use it to bind the operators and predict execution time for a set of operators instead of a single operator.

References

- [1] Pengcheng Xiong, Yun Chi, Shenghuo Zhu, Junichi Tatemura, Calton Pu, and Hakan Hacigümüş. Activesla: a profit-oriented admission control framework for database-as-a-service providers. page 15, 2011.